

Boosting

(source: David Rosenberg)

Yann Chevaleyre

November 4, 2022

1 Introduction

2 Forward Stagewise Additive Modeling (FSAM)

- Application de FSAM à la Regression: L^2 Boosting
- Application de FSAM à la Classification: Algorithme AdaBoost
- Analyser la convergence d'Adaboost: lien avec Hedge

3 Gradient Boosting / “Anyboost”

- Example: BinomialBoost

4 Gradient Tree Boosting

- GBM Regression with Stumps

Introduction

Adaptive Basis Function Model

- Base hypothesis space: \mathcal{F} of $\hat{\mathcal{Y}}$ -valued functions
- Combined hypothesis space: \mathcal{F}_M :

$$\mathcal{F}_M = \left\{ \sum_{m=1}^M v_m h_m(x) \mid v_m \in \mathbf{R}, h_m \in \mathcal{F}, m = 1, \dots, M \right\}$$

- Suppose we're given some data $S = ((x_1, y_1), \dots, (x_n, y_n))$.
- Learning is choosing $v_1, \dots, v_M \in \mathbf{R}$ and $h_1, \dots, h_M \in \mathcal{F}$ to fit S .

Note:

in bagging, we learn h_i , but $v_i = \frac{1}{M}$ for all classifiers. Boosting will learn both!!

Empirical Risk Minimization

- We'll consider learning by **empirical risk minimization**:

$$\hat{h} = \arg \min_{f \in \mathcal{F}_M} \frac{1}{n} \sum_{i=1}^n \ell(h(x_i), y_i),$$

for some **loss function** $\ell(y, \hat{y})$.

- Write ERM objective function as

$$J(v_1, \dots, v_M, h_1, \dots, h_M) = \frac{1}{n} \sum_{i=1}^n \ell \left(y_i, \sum_{m=1}^M v_m h_m(x) \right).$$

- How to optimize J ? i.e. how to learn?

- **Suppose** our base hypothesis space is parameterized by $\Theta = \mathbf{R}^d$:

$$J(v_1, \dots, v_M, \theta_1, \dots, \theta_M) = \frac{1}{n} \sum_{i=1}^n \ell \left(\sum_{m=1}^M v_m h_{\theta_m}(x), y_i \right).$$

- Can we can differentiate J w.r.t. v_m 's and θ_m 's? Optimize with SGD?
- For **some** hypothesis spaces and typical loss functions, yes!
- Neural networks fall into this category! (h_1, \dots, h_M are neurons of last hidden layer.)

What if Gradient Based Methods Don't Apply?

- What if base hypothesis space \mathcal{F} consists of decision trees?
- Can we even parameterize trees with $\Theta = \mathbb{R}^b$?
- Even if we could for some set of trees,
 - predictions would not change continuously w.r.t. $\theta \in \Theta$,
 - and so certainly not differentiable.
- Today we'll discuss **boosting**. It applies whenever
 - we can compute a particular form of the above ERM (FSAM algorithms)
 - our loss function is [sub]differentiable w.r.t. training predictions $f(x_i)$, and we can do regression with the base hypothesis space \mathcal{F} (gradient-boost).

Forward Stagewise Additive Modeling (FSAM)

Forward Stagewise Additive Modeling (FSAM)

- FSAM is an iterative optimization algorithm for fitting adaptive basis function models.
- Start with $f_0 \equiv 0$.
- After $m-1$ stages, we have

$$f_{m-1} = \sum_{i=1}^{m-1} \nu_i h_i.$$

- In m 'th round, we want to find
 - **step direction** $h_m \in \mathcal{F}$ (i.e. a basis function) and
 - **step size** $\nu_i > 0$
- such that

$$f_m = f_{m-1} + \nu_i h_m$$

improves objective function value by as much as possible.

Forward Stagewise Additive Modeling for ERM

① Initialize $f_0(x) = 0$.

② For $m = 1$ to M :

① Compute:

$$(\nu_m, h_m) = \arg \min_{\nu \in \mathbf{R}, h \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \ell \left(f_{m-1}(x_i) + \underbrace{\nu h(x_i)}_{\text{new piece}}, y_i \right).$$

② Set $f_m = f_{m-1} + \nu_m h$.

③ Return: f_M .

Application de FSAM à la Regression: L^2 Boosting

FSAM pour la Regression: L^2 Boosting

- Utilisons la “mean square error”.

$$L(v, h) = \frac{1}{n} \sum_{i=1}^N \left(y_i - \left[f_{m-1}(x_i) \quad \underbrace{+ v h(x_i)}_{\text{nouveau classifieur}} \right] \right)^2$$

- Si \mathcal{F} est “fermé par changement d’échelle” alors on peut oublier v et n’apprendre que h .
- minimiser

$$L(h) = \frac{1}{n} \sum_{i=1}^n \left(\left[\underbrace{y_i - f_{m-1}(x_i)}_{\text{résidus}} \right] - h(x_i) \right)^2$$


- Ce revient à faire un moindre carré sur les résidus !
- L’algorithme s’appelle parfois “matching pursuit”

L^2 Boosting - interprétation géométrique

(au tableau)




Regression Stumps

- A **regression stump** is a regression tree with a single split.
- A **regression stump** is a function of the form $h(x) = a1(x_i \leq c) + b1(x_i > c)$.






Plot courtesy of Brett Bernstein.

FSAM L^2 Boosting with Decision Stumps: Demo



after step 1
residue
 $= y - f_1$



Application de FSAM à la Classification: Algorithme AdaBoost


The Classification Problem

- Outcome space $\mathcal{Y} = \{-1, 1\}$
- The set of *base classifiers* is $\mathcal{F} \subset \mathcal{X} \mapsto \{-1, 1\}$ (e.g. decision stumps)
- We want to learn a *scoring function* $f_M \in \mathcal{F}_M$ where

$$\mathcal{F}_M = \left\{ \sum_{m=1}^M v_m h_m(x) \mid v_m \in \mathbf{R}, h_m \in \mathcal{F}, m = 1, \dots, M \right\}$$

- As usual, this scoring function induces a “hard”-classifier $\text{sign}(f_M(x))$
- Can we optimize a scoring loss $f_M = \arg \min_{f \in \mathcal{F}_M} \sum_{i=1}^N \ell(f(x_i), y_i)$?

Scoring Losses for Classification



- All these losses are *well calibrated*
- For these functions, a direct computation of FSAM is not easy.
- For the **exp loss**: $\ell(f(x), y) = \exp(-yf(x))$ there is an “indirect” algo: **Adaboost**

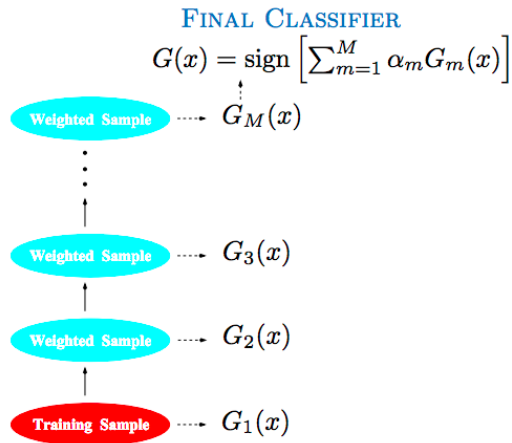
AdaBoost - Rough Sketch

- Training set $S = ((x_1, y_1), \dots, (x_n, y_n))$.
- Start with equal weight on all training points $w_1 = \dots = w_n = 1$.
- Repeat for $m = 1, \dots, M$:
 - Find base classifier $h_m(x)$ that **tries** to fit weighted training data with 0/1 loss

$$\hat{R}^w(f) = \frac{1}{W} \sum_{i=1}^N w_i \ell(f(x_i), y_i) \quad \text{where } W = \sum_{i=1}^n w_i$$

- Increase weight w_i on the points $h_m(x)$ misclassifies
- So far, we've generated M classifiers: $h_1, \dots, h_M : \mathcal{X} \rightarrow \{-1, 1\}$.
- Final scoring function is $f_M(x) = \sum_{m=1}^M v_m h_m(x)$, for some weights v_m .

AdaBoost: Schematic



From ESL Figure 10.1

AdaBoost: Algorithm

Given training set $S = \{(x_1, y_1), \dots, (x_n, y_n)\}$.


- ① Initialize observation weights $w_i = 1, i = 1, 2, \dots, N$.
- ② For $m = 1$ to M :
 - ① learner fits weighted training data and returns $h_m(x)$
 - ② Compute **weighted empirical 0-1 risk**:

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i 1(y_i \neq h_m(x_i)) \quad \text{where } W = \sum_{i=1}^n w_i.$$


- ③ Compute $v_m = \ln\left(\frac{1-\text{err}_m}{\text{err}_m}\right)$ [**classifier weight**]
 - ④ Set $w_i \leftarrow w_i \cdot \exp[v_m 1(y_i \neq h_m(x_i))]$, $i = 1, 2, \dots, n$ [**example weight adjustment**]
- ③ Output $f_M(x) = \sum_{m=1}^M v_m h_m(x)$.

AdaBoost: Pas à pas

AdaBoost Classifier Working Principle with Decision Stump as a Base Classifier



AdaBoost: Pas à pas



Adaboost en tant qu'algorithme FSAM

- Soit $\ell(\hat{y}, y) = \exp(-\hat{y}y)$. Montrons qu'à chaque étape d'Adaboost, l'algorithme calcule

$$\begin{aligned}(\nu_m, h_m) &= \arg \min_{\nu \in \mathbf{R}, h \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \ell \left(f_{m-1}(x_i) + \underbrace{\nu h(x_i)}_{\text{new piece}}, y_i \right) \\&= \arg \min_{\nu} \frac{1}{n} \sum_i \exp(-y_i (f_{m-1}(x_i) + \nu h(x_i))) \\&= \arg \min_{\nu} \frac{1}{n} \sum_i \alpha_i \exp(-y_i \nu h(x_i)) \\&\quad \text{with } \alpha_i = \exp(-y_i f_{m-1}(x_i))\end{aligned}$$

Analyser la convergence d'Adaboost: lien avec Hedge

Simplified variant of AdaBoost

- In this section, we will analyse a simplified version of Adaboost, which we will relate to the online learning *Hedge* algorithm.

Given training set $S = \{(x_1, y_1), \dots, (x_n, y_n)\}$.

- 1 Initialize observation weights $w_i = 1$, $i = 1, 2, \dots, N$.
- 2 For $m = 1$ to M :
 - 1 learner fits weighted training data and returns $h_m(x)$
 - 2 Compute **weighted empirical 0-1 risk**:

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i 1(y_i \neq h_m(x_i))$$

- 3 Set $w_i \leftarrow w_i \cdot \exp[-\beta \cdot 1(y_i \neq h_m(x_i))]$, $i = 1, 2, \dots, n$ [**example weight adjustment**]
- 3 Output $f_M(x) = \sum_{m=1}^M h_m(x)$.

Rappel: apprentissage en-ligne en 0/1 loss

Protocole

Pour $t = 1$ à T

- L'environnement choisit x_t, y_t , et révèle x_t à l'apprenant
- **L'apprenant prédit \hat{y}_t (en général, il choisit h_t et prédit $\hat{y}_t = h_t(x_t)$)**
- L'environnement révèle y_t
- L'apprenant reçoit le cout de $\ell^{0/1}(\hat{y}_t, y_t)$

Objectif: produire la séquence de classifieurs $h_1 \dots h_T$ tels que la perte cumulée $\sum_{t=1}^T \ell^{0/1}(h_t(x_t), y_t)$ soit minimisée

Rappel: Hedge

- Je choisis $P_t(h) = \frac{1}{\Omega_t} w_{h,t}$ avec
 - $w_{h,1} = 1$
 - $w_{h,t+1} = w_{h,t} e^{-\beta \ell^{0/1}(h(x_t), y_t)}$ pour une constante $\beta > 0$

Algorithme Hedge (voir cours d'apprentissage en ligne)

$\mathcal{F}_1 = \mathcal{F}$

Pour $t = 1$ à T

- je reçois x_t
- **je tire** $h_t \sim P_t$
- je reçois le vrai label y_t , et ma prédiction me coute $\ell(h_t(x_t), y_t)$
- **je mets à jour** P_{t+1}

Thm: in expectation, $\text{Regret}_T = \sum_{t=1}^T \ell_t(h_t) - \min_{h \in \mathcal{F}} \sum_{t=1}^T \ell_t(h) \leq \sqrt{2T \ln |\mathcal{F}|}$

Problème dual de l'apprentissage en-ligne

- (on inverse la place de l'apprenant et de l'environnement)

Protocole dual

Pour $t = 1$ à T

- L'environnement choisit sans le révéler un classifieur h_t
- **L'apprenant choisit x_t, y_t de l'ensemble S**
- L'environnement révèle h_t
- L'apprenant reçoit le **gain** de $\ell(h_t(x_t), y_t)$

Objectif: *Trouver les exemples $(x_t, y_t) \in S$ qui maximisent la perte cumulée $\sum_{t=1}^T \ell^{0/1}(h_t(x_t), y_t)$, ou qui minimisent $\sum_{t=1}^T (1 - \ell^{0/1}(h_t(x_t), y_t))$*

Dual Hedge

- Je choisis $P_t(i) = \frac{1}{\Omega_t} w_{i,t}$, distribution discrète sur $S = \{(x_1, y_1) \dots (x_N, y_N)\}$ avec:
 - $w_{i,1} = 1$, et $w_{i,t+1} = w_{i,t} e^{-\beta(1-\ell(h_t(x_i), y_i))}$ pour une constante $\beta > 0$

Algorithme Dual-Hedge

Pour $t = 1$ à T

- l'environnement choisit h_t sans le révéler
- **je tire** $(x_t, y_t) \sim P_t$
- je reçois h_t et mon choix me coute $1 - \ell(h_t(x_t), y_t)$
- **je mets à jour** P_{t+1} en calculant $w_{i,t+1} = w_{i,t} e^{-\beta(1-\ell(h_t(x_i), y_i))}$ pour tout i

Thm: we have in expectation,

$$\text{Regret}_T = \sum_{t=1}^T (1 - \ell(h_t(x_t), y_t)) - \min_{i \in \{1..N\}} \sum_{t=1}^T (1 - \ell(h_t(x_t), y_t)) \leq \sqrt{2T \ln |S|}$$

Dual Hedge

- La borne de regret est vraie pour TOUTE stratégie de l'environnement. On peut donc imposer une condition sur l'environnement sans changer la borne.

Algorithme Dual-Hedge

Pour $t = 1$ à T

- l'environnement choisit h_t tel que $err_t = \frac{1}{W} \sum_{i=1}^N w_i 1(y_i \neq h_t(x_i)) \leq \frac{1}{2} - \gamma$
- ...
- **je mets à jour** P_{t+1} en calculant $w_{i,t+1} = w_{i,t} e^{-\beta(1-\ell(h_t(x_i), y_i))}$ pour tout i

Thm:

$$\begin{aligned} \text{Regret}_T &= \mathbb{E} \left[\sum_{t=1}^T (1 - \ell(h_t(x_t), y_t)) - \min_{i \in \{1..N\}} \sum_{t=1}^T (1 - \ell(h_t(x_t), y_t)) \right] \\ &= \sum_{t=1}^T \mathbb{E}_{x,y \sim P_t} \left[(1 - \ell(h_t(x), y)) - \min_{i \in \{1..N\}} \sum_{t=1}^T (1 - \ell(h_t(x), y)) \right] \leq \sqrt{2T \ln |S|} \end{aligned}$$

Analyse de Dual-Hedge / Lien avec (Simplified) Adaboost

- Nous allons montrer que si, à chaque étape, h_t satisfait

$$err_t = \frac{1}{W} \sum_{i=1}^N w_i 1(y_i \neq h_t(x_i)) \leq \frac{1}{2} - \gamma \quad (\text{weak learning hypothesis})$$

pour $\gamma \in]0, \frac{1}{2}[$, alors au bout de quelques étapes, le classifieur final aura une erreur empirique nulle.

- Notons que $err_t = \mathbb{E}_{x,y \sim P_t} (\ell^{0/1}(h_t(x), y))$.
- Plus précisément:

Thm de convergence d'Adaboost: Sous l'hypothèse de weak learning, après $T = \frac{2}{\gamma^2} \ln N$ pas de temps, le classifieur majoritaire a un taux d'erreur de classification nulle sur l'échantillon S . (preuve au tableau)

Gradient Boosting / “Anyboost”

FSAM Is Iterative Optimization

- The FSAM step

$$(\mathbf{v}_m, h_m) = \arg \min_{\mathbf{v} \in \mathbf{R}, h \in \mathcal{F}} \sum_{i=1}^n \ell \left(y_i, f_{m-1}(x_i) + \underbrace{\mathbf{v} h(x_i)}_{\text{new piece}} \right).$$

- Hard part: finding the **best step direction** h .
- What if we looked for the **locally best** step direction?
 - like in gradient descent

“Functional” Gradient Descent

- We want to minimize

$$J(f) = \sum_{i=1}^n \ell(y_i, f(x_i)).$$

- In some sense, we want to take the gradient w.r.t. “ f ”, whatever that means.
- $J(f)$ only depends on f at the n training points.
- Define

$$\mathbf{f} = (f(x_1), \dots, f(x_n))^T$$

and write the objective function as

$$J(\mathbf{f}) = \sum_{i=1}^n \ell(y_i, \mathbf{f}_i).$$

Functional Gradient Descent: Unconstrained Step Direction

- Consider gradient descent on

$$\mathbf{f}_i = f(x_i) \\ J(\mathbf{f}) = \sum \ell(y_i, \mathbf{f}_i)$$

$$J(\mathbf{f}) = \sum_{i=1}^n \ell(y_i, \mathbf{f}_i).$$

$$\mathbf{f} = \begin{pmatrix} f_1 \\ \vdots \\ f_n \end{pmatrix} \in \mathbb{R}^n \\ \mathbf{g} \in \mathbb{R}^n$$

- The **negative gradient step direction** at \mathbf{f} is

$$\begin{aligned} -\mathbf{g} &= -\nabla_{\mathbf{f}} J(\mathbf{f}) \\ &= -(\partial_{\mathbf{f}_1} \ell(y_1, \mathbf{f}_1), \dots, \partial_{\mathbf{f}_n} \ell(y_n, \mathbf{f}_n)) \end{aligned}$$

which we can easily calculate.

- $-\mathbf{g} \in \mathbb{R}^n$ is the direction we want to change each of our n predictions on training data.
- Eventually we need more than just \mathbf{f} , we'll need the function f .

Functional Gradient Descent: Projection Step

- Unconstrained step direction is

$$-\mathbf{g} = -\nabla_{\mathbf{f}} J(\mathbf{f}) = -(\partial_{\mathbf{f}_1} \ell(y_1, \mathbf{f}_1), \dots, \partial_{\mathbf{f}_n} \ell(y_n, \mathbf{f}_n)).$$

- Also called the “**pseudo-residuals**”
 - (for square loss, they’re exactly the residuals)
- Find the closest base hypothesis $h \in \mathcal{F}$ (in the ℓ^2 sense):

$$\sum \ell(y_i, f_i - g_i) \leq \sum \ell(y_i, f_i)$$

$$\min_{h \in \mathcal{F}} \sum_{i=1}^n (-\mathbf{g}_i - h(x_i))^2.$$

$$\ell(y_i, f_{m-1} + h(x_i))$$

- This is a least squares regression problem over hypothesis space \mathcal{F} .
- Take the $h \in \mathcal{F}$ that best approximates $-\mathbf{g}$ as our step direction.

Functional Gradient Descent: Step Size

- Finally, we choose a stepsize.
- Option 1 (Line search):

$$\nu_m = \arg \min_{\nu > 0} \sum_{i=1}^n \ell\{y_i, f_{m-1}(x_i) + \nu h_m(x_i)\}.$$

- Option 2: (learning rate parameter – **more common**)
 - We consider $\nu = 1$ to be the full gradient step.
 - Choose a fixed $\nu \in (0, 1)$ – called a **learning rate or shrinkage parameter**.
 - A value of $\nu = 0.1$ is typical – optimize as a hyperparameter .

The Gradient Boosting Machine Ingredients (Recap)

- Take any loss function [sub]differentiable w.r.t. the prediction
- Choose a base hypothesis space for regression.
- Choose number of steps (or a stopping criterion).
- Choose step size methodology.
- Then you're good to go!

Example: BinomialBoost

BinomialBoost: Gradient Boosting with Logistic Loss

- Recall the logistic loss for classification, with $\mathcal{Y} = \{-1, 1\}$:

$$\ell(y, f(x)) = \log(1 + e^{-yf(x)})$$

- Pseudoresidual for i 'th example is negative derivative of loss w.r.t. prediction:

$$\begin{aligned} r_i &= -\partial_{f(x_i)} \left[\log(1 + e^{-y_i f(x_i)}) \right] \\ &= \frac{y_i e^{-y_i f(x_i)}}{1 + e^{-y_i f(x_i)}} \\ &= \frac{y_i}{1 + e^{y_i f(x_i)}} \end{aligned}$$

BinomialBoost: Gradient Boosting with Logistic Loss

- Pseudoresidual for i th example:

$$r_i = -\partial_{f(x_i)} \left[\log \left(1 + e^{-y_i f(x_i)} \right) \right] = \frac{y_i}{1 + e^{y_i f(x_i)}}$$

- So if $f_{m-1}(x)$ is prediction after $m-1$ rounds, step direction for m 'th round is

$$h_m = \arg \min_{h \in \mathcal{F}} \sum_{i=1}^n \left[\left(\frac{y_i}{1 + e^{y_i f_{m-1}(x_i)}} \right) - h(x_i) \right]^2.$$

- And $f_m(x) = f_{m-1}(x) + \eta h_m(x)$.

Gradient Tree Boosting

Gradient Tree Boosting

- One common form of gradient boosting machine takes


$$\mathcal{F} = \{\text{regression trees of size } J\},$$

where J is the number of terminal nodes.

- $J = 2$ gives decision stumps
- HTF recommends $4 \leq J \leq 8$ (but more recent results use much larger trees)
- Software packages:
 - Gradient tree boosting is implemented by the **gbm package** for R
 - as `GradientBoostingClassifier` and `GradientBoostingRegressor` in **sklearn**
 - **xgboost** and **lightGBM** are state of the art for speed and performance


GBM Regression with Stumps

Sinc Function: Our Dataset



From Natekin and Knoll's "Gradient boosting machines, a tutorial"

Minimizing Square Loss with Ensemble of Decision Stumps



Decision stumps with 1, 10, 50, and 100 steps, step size $\lambda = 1$.

Figure 3 from Natekin and Knoll's "Gradient boosting machines, a tutorial"

Rule of Thumb

- The smaller the step size, the more steps you'll need.
- But never seems to make results worse, and often better.
- So set your step size as small as you have patience for.