

Neural Nets for text classification

Alexandre Allauzen

Fall 2023



Roadmap

Call for Participation

Introduction

Word embeddings

Convolution for text

Outline

Call for Participation

Introduction

Word embeddings

Convolution for text

Who wants to do a presentation on:

- ELECTRA
- Mega
- State-Space model
- RWKV
- Pegasus
- LORA
- Adapter, ...

Schedules

- 10 minutes presentation
- a draft version of the slides to have feedback: 14/11
- presentation: 21/11

Outline

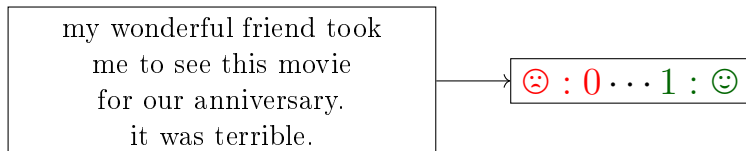
Call for Participation

Introduction

Word embeddings

Convolution for text

Text classification/rating



How to represent the input text ?

- Bag of features (words, ...)
- Really represent the sentence as a whole

Bag of words (BOW)

this movie is just great , with a great music , while a bit long

vocabulary	binary bag	count bag	tf.idf bag	...
awesome	0	0	0	...
great	1	2	1.9	...
long	1	1	2.5	...
the	0	0	0	...
this	1	1	0.1	...

A basic vectorial representation of text

$$\mathbf{x} = \begin{pmatrix} 0 \\ 2 \\ 1 \\ 0 \\ 1 \end{pmatrix} \in \mathbb{R}^D$$

$$\left. \begin{array}{l} \textit{awesome} \\ \textit{great} \\ \textit{long} \\ \textit{the} \\ \textit{this} \end{array} \right\} D$$

A simple problem

Assumptions

- Let define a finite set of known words: the vocabulary \mathcal{V}
- A text is a vector \mathbf{x} of dimension $D = |\mathcal{V}|$
- Each component encodes the presence of a word

Then machine learning

- Naive Bayes
- SVM, Random Forrest, ...
- **Logistic Regression**



Logistic regression

The class c is the outcome of the binary random variable C

The sigmoid/logistic function

$$a = w_0 + \mathbf{w}^t \mathbf{x} \in \mathbb{R}$$

$$\sigma(a) = \frac{e^a}{1 + e^a} = \frac{1}{1 + e^{-a}} \text{ and } y = P(C = 1|\mathbf{x}) = \sigma(w_0 + \mathbf{w}^t \mathbf{x})$$



Training a Logistic regression model

- The parameters are $\boldsymbol{\theta} = (w_0, \mathbf{w})$,
- The i.i.d dataset: $\mathcal{D} = (\mathbf{x}_{(i)}, c_{(i)})_{i=1}^n$

Loss function minimization

$$\begin{aligned}\mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) &= - \sum_{i=1}^n \log(P(C = c_{(i)} | \mathbf{x}; \boldsymbol{\theta})) \\ &= - \sum_{i=1}^n (c_{(i)} \log(y_{(i)}) + (1 - c_{(i)}) \log(1 - y_{(i)})) \\ y_{(i)} &= \sigma(w_0 + \mathbf{w}^t \mathbf{x}_{(i)})\end{aligned}$$

Optimization method

Stochastic Gradient Descent, or improved version (ADAM, L-BFGS, ...)

Outline

Call for Participation

Introduction

Word embeddings

Convolution for text

Back to logistic regression

$$\mathbf{x} = \begin{pmatrix} 0 \\ 2 \\ 1 \\ 0 \\ 1 \end{pmatrix} \in \mathbb{R}^D \quad \left. \begin{array}{l} \textit{awesome} \\ \textit{great} \\ \textit{long} \\ \textit{the} \\ \textit{this} \end{array} \right\} D$$

For one input text:


$$w_0 + \mathbf{w}^t \mathbf{x} = w_0 + \textcolor{red}{2} \times w_2 + w_3 + w_5$$

The class is positive ($y = 1$) if

$$\begin{aligned} w_0 + 2 \times w_2 + w_3 + w_5 &> 0 \\ 2 \times w_{\textit{great}} + w_{\textit{long}} + w_{\textit{this}} + &> -w_0 \end{aligned}$$

A limited representation of words

With the logistic regression model on a bag of words:



Consider the two following examples:

the end is really bad		☹️ $\Rightarrow w_{\text{bad}}$	\searrow
the bad guy is <i>awesome</i>		😊 $\Rightarrow w_{\text{bad}}$	$\searrow, w_{\text{awesome}} \nearrow$

Multiple dimensions could help to:

- represent different usage
- consider the context
- leverage more from sparse, sometime ambiguous observations.

A simple model for document classification - part 1

Idea

- The word representation could be shared among classes
- While their interpretation depends on the class

Input representation and composition

$$\mathbf{R} \times \mathbf{x} = \begin{pmatrix} \vdots & \vdots & \vdots & \vdots & \vdots \\ \mathbf{v}_1 & \mathbf{v}_2 & \mathbf{v}_3 & \mathbf{v}_4 & \mathbf{v}_5 \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{pmatrix} \times \begin{pmatrix} 0 \\ \mathbf{2} \\ \mathbf{1} \\ 0 \\ \mathbf{1} \end{pmatrix} = 2 \times \mathbf{v}_2 + \mathbf{v}_3 + \mathbf{v}_5 = \mathbf{d}$$

A simple model for document classification - part 2

Classification

$$\begin{aligned}P(y|\mathbf{x}) &= \text{softmax}(\mathbf{W}^o \mathbf{d}) = \text{softmax}(\mathbf{W}^o \times \mathbf{R}\mathbf{x}), \text{ or} \\ &= \text{softmax}(\mathbf{W}^o \times f(\mathbf{R}\mathbf{x})),\end{aligned}$$

with f a non-linear activation function.

Parameters


$$\theta = (\mathbf{R}, \mathbf{W}^o) \rightarrow \text{to learn !!}$$

Reminder

If $\mathbf{y} = \text{softmax}(\mathbf{a})$, \mathbf{y} is a vector and \mathbf{a} is called the logit vector

$$y_i = \frac{e^{a_i}}{\sum_j e^{a_j}}$$

A first neural network



- $\mathbf{x} : (|\mathcal{V}|, 1)$
- $\mathbf{R} : (K, |\mathcal{V}|)$
- $\mathbf{d} : (K, 1)$
- $\mathbf{W}^o : (1, K)$
- $y : (1, 1)$

$$\mathbf{d} = \mathbf{R} \times \mathbf{x}$$

$$y = \sigma(\mathbf{W}^o \times \mathbf{d})$$

Word embeddings

Definitions:

- To each word, a continuous vector is associated: its **embedding**.
- The matrix \mathbf{R} is called the **look-up table** and store the word embeddings.

Note:

- The term *look-up* comes from the real operation $\mathbf{R} \times \mathbf{x}$ is only theoretical !
- No computational cost, only storage and trainability challenge (enough observations for each word, Zipf, ...)
- **Pre-training** and **fine-tuning**

Unsupervised Pre-training of Word Embeddings

The question

- How to efficiently learn word representation
- based on the observation of raw texts ?

Distributional representations

You shall know a word by the company it keeps (Firth, J. R., 1957)

and

Words are similar if they appear in similar contexts (Harris 1954).

In practice

Word2Vec [6]

Context Bag of Words (CBOW)

The game

southern trees [??] strange fruits

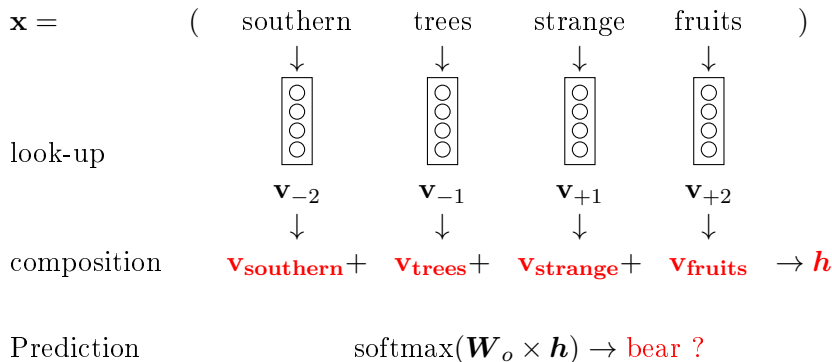
Guess the word in the middle !

Context Bag of Words (CBOW)

The game

southern trees [??] strange fruits

Guess the word in the middle !



CBOW: details

Fast pre-training of word embeddings

- Introduced in [6] as a simplification of [1] (neural language model)
- Trained with negative sampling (Closed to Noise Contrastive Estimation [2])
- An efficient and tractable approximation of the count based method [5]

Other flavor

- Skip-gram [6]
- Glove [7]
- Fasttext [3]

CBOW: Maximum Likelihood Estimate

In $P(w|\mathbf{x}; \boldsymbol{\theta})$:

- predict the word w in the middle,
- given \mathbf{x} the context.

MLE

$$\mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) = - \sum_{i=1}^n \log(P(C = w|\mathbf{x}; \boldsymbol{\theta})),$$

- The probability distribution over \mathcal{V} is given by a softmax
- The set of possible outcomes is \mathcal{V} .

Cost of the softmax

$$\mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) = - \sum_{(\mathbf{x}, \hat{w}) \in \mathcal{D}} \log P_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x})$$

$$P_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x}) = \frac{e^{s_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x})}}{\sum_{w' \in \mathcal{V}} e^{s_{\boldsymbol{\theta}}(w'|\mathbf{x})}}$$

$$\log P_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x}) = s_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x}) - \log \left(\sum_{w' \in \mathcal{V}} e^{s_{\boldsymbol{\theta}}(w'|\mathbf{x})} \right)$$

$$\frac{\partial \log P_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x})}{\partial \boldsymbol{\theta}} = \frac{\partial s_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x})}{\partial \boldsymbol{\theta}} - \underbrace{\sum_{w' \in \mathcal{V}} P_{\boldsymbol{\theta}}(w'|\mathbf{x}) \frac{\partial s_{\boldsymbol{\theta}}(w', \mathbf{x})}{\partial \boldsymbol{\theta}}}_{\text{costly!}}$$

Negative sampling

Recast the problem as a binary classification task:

- Positive examples: $(\mathbf{x}, w) \in \mathcal{D}$
- Negative examples: (\mathbf{x}, \tilde{w}) , with $\tilde{w} \sim \mathcal{V}$

Use a binary classifier !

In practice:

- for one positive example $\sim \mathcal{D}$
- sample K negative and random samples from \mathcal{V}
- K is small (compared to the size of \mathcal{V})
- the noise distribution does matter

Outline

Call for Participation

Introduction

Word embeddings

Convolution for text

Local contexts

the	end	is	very	bad	but	what	a	great	music
-----	-----	----	------	-----	-----	------	---	-------	-------

Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$ $very \rightarrow bad ++$						

Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$ $very \rightarrow bad++$						
			$\underbrace{\hspace{2.5cm}}$ $but \text{ will change } bad$						

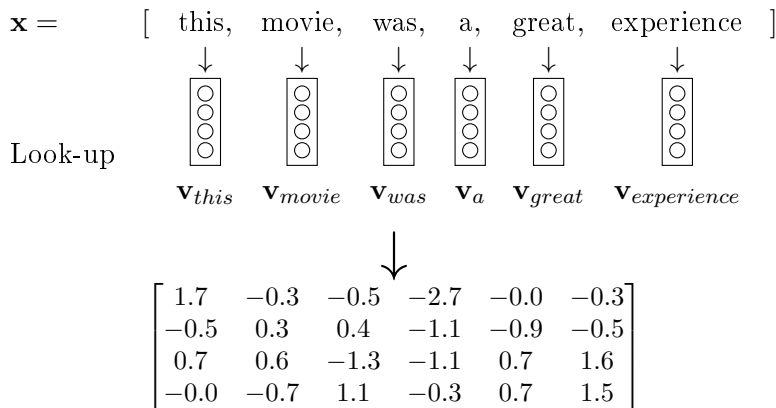
Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$ <i>very</i> \rightarrow <i>bad</i> ++						
			$\underbrace{\hspace{2.5cm}}$ <i>but</i> will change <i>bad</i>						
		$\underbrace{\hspace{3.5cm}}$ <i>bad</i> is for <i>end</i> not <i>music</i>						$\underbrace{\hspace{1.5cm}}$ <i>great</i> is for <i>music</i> not fo <i>end</i>	

Motivations

- Local contextualisation
- Global view of the sentence

Another view of a sentence



Convolutional Neural Networks for Sentence Classification

- A short paper of 2014 [4]
- A simple and SOTA paper on text classification

Convolution in 1D

Extract a frame, or a window, and apply a “filter”

The filter

Kernel size of $ks = 2$

$w_{1,1}$	$w_{1,2}$
$w_{2,1}$	$w_{2,2}$
$w_{3,1}$	$w_{3,2}$
$w_{4,1}$	$w_{4,2}$



The input sequence

$L = 6$ vectors in \mathbb{R}^D , $D = 4$

$x_{1,1}$	$x_{1,2}$	$x_{1,3}$	$x_{1,4}$	$x_{1,5}$	$x_{1,6}$
$x_{2,1}$	$x_{2,2}$	$x_{2,3}$	$x_{2,4}$	$x_{2,5}$	$x_{2,6}$
$x_{3,1}$	$x_{3,2}$	$x_{3,3}$	$x_{3,4}$	$x_{3,5}$	$x_{3,6}$
$x_{4,1}$	$x_{4,2}$	$x_{4,3}$	$x_{4,4}$	$x_{4,5}$	$x_{4,6}$

The output value (output channel)

$$\text{At time } t = 1, \text{ } h_1 = \sum_{i,j} w_{i,j} \times x_{i,j}$$

Convolution in 1D : time $t = 2$

Stride = 1

The filter

Kernel size of $ks = 2$

$w_{1,1}$	$w_{1,2}$
$w_{2,1}$	$w_{2,2}$
$w_{3,1}$	$w_{3,2}$
$w_{4,1}$	$w_{4,2}$



The input sequence

$L = 6$ vectors in \mathbb{R}^D , $D = 4$

$x_{1,1}$	$x_{1,2}$	$x_{1,3}$	$x_{1,4}$	$x_{1,5}$	$x_{1,6}$
$x_{2,1}$	$x_{2,2}$	$x_{2,3}$	$x_{2,4}$	$x_{2,5}$	$x_{2,6}$
$x_{3,1}$	$x_{3,2}$	$x_{3,3}$	$x_{3,4}$	$x_{3,5}$	$x_{3,6}$
$x_{4,1}$	$x_{4,2}$	$x_{4,3}$	$x_{4,4}$	$x_{4,5}$	$x_{4,6}$

The output value (output channel)

At time $t = 2$, $h_2 = \sum_{i,j} w_{i,j} \times x_{i+1,j}$

Convolution in 1D : time $t = 5$

Stride = 1

The filter

Kernel size of $ks = 2$

The input sequence

$L = 6$ vectors in \mathbb{R}^D , $D = 4$

$w_{1,1}$	$w_{1,2}$
$w_{2,1}$	$w_{2,2}$
$w_{3,1}$	$w_{3,2}$
$w_{4,1}$	$w_{4,2}$




$x_{1,1}$	$x_{1,2}$	$x_{1,3}$	$x_{1,4}$	$x_{1,5}$	$x_{1,6}$
$x_{2,1}$	$x_{2,2}$	$x_{2,3}$	$x_{2,4}$	$x_{2,5}$	$x_{2,6}$
$x_{3,1}$	$x_{3,2}$	$x_{3,3}$	$x_{3,4}$	$x_{3,5}$	$x_{3,6}$
$x_{4,1}$	$x_{4,2}$	$x_{4,3}$	$x_{4,4}$	$x_{4,5}$	$x_{4,6}$

The output value (output channel)

$$\text{At time } t = 5, \text{ } h_5 = \sum_{i,j} w_{i,j} \times x_{i+5,j}$$

Feature extraction for subsequences

Embeddings



After the convolution

$$\mathbf{h} = (\underbrace{h_1}_{(\text{this, movie})}, \underbrace{h_2}_{(\text{movie, was})}, \underbrace{h_3}_{(\text{was, a})}, \underbrace{h_4}_{(\text{a, great})}, \underbrace{h_5}_{(\text{great, experience})})$$

Convolution 1D on a embedding sequence

Local features extraction

- Each Kernel / Filter application gives **one local feature**.
- The feature relates to a n -gram.
- The **kernel size** defines the scope of each feature.
- The **stride** controls the amount of slides.

A sequence of vector \rightarrow a sequence of features


More features ?

- More Filters !
- More "output channels"

Convolution with two output channels

Output channels = 2

Input channels = 1




The output value (output channel)

$$h_{1,1} = \sum_{i,j} w_{i,j}^{(1)} \times x_{i,j}$$

$$h_{2,1} = \sum_{i,j} w_{i,j}^{(2)} \times x_{i,j}$$


Another view for two output channels



- Two filters applied to the same frame (or window)
- Two projections
- \mathbf{W} : the parameters of the filters
- \mathbf{W} is learnt

Pool !

Compress the "local" information along one dimension (e.g time)




- Mean pooling
- Max pooling, and k -max pooling

For example with:

$$\begin{bmatrix} 1.7 & -0.3 & -0.5 & -2.7 & -0.0 & -0.3 \\ -0.5 & 0.3 & 0.4 & -1.1 & -0.9 & -0.5 \end{bmatrix} \rightarrow \begin{bmatrix} -0.3 \\ -0.4 \end{bmatrix} \text{ or } \begin{bmatrix} 1.7 \\ 0.4 \end{bmatrix} \text{ or } \dots$$

Pooling can also apply to sliding windows

More Convolutions and pooling



And they can be combined (concatenation).

Convolutional Neural Networks for Sentence Classification

A summary of [4]

- Window (kernel) sizes : 3, 4, 5 with 100 feature maps for each
- Static/non-static/random/multi-channel word embeddings
- Auxiliary data for word embeddings: $\tilde{w2v}$ trained on 100 billion words from Google News ($dim = 300$)
- dropout on the penultimate layer (after the max-pooling)
- Relu and early stopping

CNN applications for NLP

Word level

- The unit = a word (the vocabulary ?)
- Compose word representation to derive a sentence representation
- Extract n -gram patterns (phrasal)

Char level

- The unit = the char (closed vocab)
- Compose chars to infer a word representation
- Extract morphological features (mostly concatenative)
unbelievable, untractable, believer, writter, forever, ...